

- 4.1 **Logistic regression:** we consider data on wage of professors in the United States over nine month periods. The profsalary dataset contains information about the participants.
- sex: binary, either man (0) or woman (1);
 - rank: categorical, one of assistant (1), associate (2) or full professor (3);
 - degree: highest degree, either masters (0) or doctorate (1);
 - yd: number of years since last degree;
 - yr: number of years in academic rank;
 - salary: salary in USD over nine months;
- (a) Fit a logistic regression to model the probability that a professor has a salary superior to 105K USD as a function of degree, sex, yr and yd. Write the equation for the mean and interpret the estimated coefficients of the model.
- (b) If you add the covariate rank, what happens? Do you identify any problem with the model? If so, try to find an explanation.
- 4.2 An education researcher is interested in the association between the number of awards students receive at a high school as a function of their math scores and the type of school they attend. The awards data contains
- awards: response variable indicating the number of prize received throughout the year
 - math: score of students on their math final exam
 - prog: student program in which the student, one of general (1), academic (2) and vocational (3).
- Fit a Poisson model and a negative binomial model with math and prog as covariates and interpret the parameters. Compare the results between the two and say which model is more appropriate, if any.
- 4.3 **Rate data** The ceb data contains information about the number of children ever born (CEB) from the Fiji Fertility Survey. The variable measured for each group of women are
- nwom: number of women in the group.
 - nceb: response, number of children ever born.
 - dur: time (in years) since wedding, either 0–4 (1), 5–9 (2), 10–14 (3), 15–19 (4), 20–24 (5) and greater than 25 (6).
 - res: categorical variable for residence, one of Suva (1), urban (2) or rural (3).
 - educ: ordinal variable giving the educational achievements, one of none (1), lower primary (2), upper primary (3), high school or higher (4).
 - var: estimated within-group variance in number of children ever born per group.
- (a) Plot (a) the number of children ever born (nceb) as a function of the number of women in the group (nwom) and (b) the mean number of children ever born per group against the variance of the number of children ever born. Comment on the two plots.
- (b) Should an offset term be included? Explain.
- If no offset is used, which function (if any) of nwom should be included in the mean model?
 - If one considers using an offset, how does it compare relative to the model with $\log(\text{nwom})$?
- (c) Fit a Poisson regression model with an offset including the three categorical covariates dur, res and educ as main effects. Which of the three predictors is the most significant?
- (d) Interpret the coefficients of the fitted model.
- (e) Assess whether there is need for an interaction between educ and dur by performing a likelihood ratio test.
- (f) It is possible to assess goodness-of-fit using diagnostic plots for the so-called deviance residuals from a Poisson generalized linear model.¹ Using software, produce diagnostic plots of (a) fitted values against deviance-based residuals, (b) quantile-quantile plot of deviance-based residuals, (c) leverage and (d) Cook distance against observation number. Comment on the adequacy of the model with all three categorical covariates and an offset. *To produce the plots in SAS, use the options*

¹In general, however, these diagnostics are harder to interpret because the observations are discrete whereas the fitted mean is continuous.

```
proc genmod data=... plots=(resdev(xbeta) leverage cooks)
```

The quantile-quantile plot can be produced with the procedure `univariate` using the standardized deviance residuals (`stdresdev`). In R, use the function `boot::glm.diag.plots` to produce graphical diagnostics.

4.4 Understanding the drivers of BIXI rentals: BIXI is a Montreal-based bicycle rental company. We examine the data for 500 days during the period 2017–2019 at the Edouard Montpetit bike docking station in front of HEC. Our interest is in explaining variability in daily bike rental (measured through the number of users) at that station based on time of the week and meteorological factors. The data consist of

- `nusers`: number of daily users at the station.
- `temp`: temperature (in degree Celcius)
- `relhumid`: percentage of relative humidity, taking values between 0 and 100.
- `weekday`: categorical variable for week day, between Sunday (1) and Saturday (7).
- `weekend`: binary variable taking value zero if rental is on a weekend (Saturday or Sunday) and one otherwise.

We consider four competing models for the data

- Model 4.4.1 is a Poisson regression model with `weekend` as covariate.
 - Model 4.4.2 is a Poisson regression model with `weekend`, `relhumid` and `temp` as covariates.
 - Model 4.4.3 is a negative binomial model with `weekend`, `relhumid` and `temp` as covariates.
 - Model 4.4.4 is a negative binomial model with `weekday` (categorical), `relhumid` and `temp` as covariates.
- (a) Is Model 4.4.1 an adequate simplification of Model 4.4.2? Assess this hypothesis formally.
- (b) Interpret the coefficients for the intercept and for `relhumid` in Model 4.4.2.
- (c) Compare the model fit of the negative binomial model (Model 4.4.3) with that of the Poisson (Model 4.4.2) using all of the following methods: (a) the deviance statistic (b) a likelihood ratio test and (c) information criteria.
- (d) Suppose that, rather than including `weekend`, we instead consider `weekday` as covariate in the models. Explain how the model would differ if we included `weekday` as an integer-valued variable as opposed to declaring it categorical. Which of the two makes more sense in the present context?
- (e) The equation for the mean of Model 4.4.4 is

$$E(\text{nusers}) = \exp(\beta_0 + \beta_1 \text{temp} + \beta_2 \text{relhumid} + \beta_3 \mathbf{1}_{\text{weekday}=2} + \beta_4 \mathbf{1}_{\text{weekday}=3} + \beta_5 \mathbf{1}_{\text{weekday}=4} + \beta_6 \mathbf{1}_{\text{weekday}=5} + \beta_7 \mathbf{1}_{\text{weekday}=6} + \beta_8 \mathbf{1}_{\text{weekday}=7}).$$

Write the null hypothesis for the test comparing Model 4.4.3 to Model 4.4.4 in terms of the model parameters $\boldsymbol{\beta}$, thereby showing that Model 4.4.3 is nested within Model 4.4.4. Does the number of user significantly vary between weekdays and between weekend days?

4.5 Soccer matches: Let Y_{ij} (resp. Z_{ij}) denote the score of the home (resp. visitor) team for a soccer match opposing teams i and j . Maher (1982) suggested modelling the scores as

$$Y_{ij1} \sim \text{Po}\{\exp(\delta + \alpha_i + \beta_j)\}, \quad Y_{ij2} \sim \text{Po}\{\exp(\alpha_j + \beta_i)\}, \quad i \neq j; i, j \in \{1, \dots, 20\}, \quad (\text{E4.5.1})$$

where α_i represent the offensive strength of the team, β_j the defensive strength of team j and δ is the common home advantage. The scores in a match and between matches are assumed to be independent of one another. The data set `soccer` contains the results of football (soccer) matches for the 2015 season of the English Premier Ligue (EPL) and contains the following variables

- `score`: number of goals of team during a match
- `team`: categorical variable giving the name of the team which scored the goals
- `opponent`: categorical variable giving the name of the adversary
- `home`: binary variable, 1 if `team` is playing at home, 0 otherwise.

- (a) A common home advantage δ makes sense provided that the scores at home and away are independent, i.e., there is no interaction between the two. To validate this hypothesis, we consider aggregates over multiple seasons of the scores, cross-classified in terms of number of points for the team at home and the team away, for each match (Table 1). The file socceragg contains the two-way contingency table in long-format. Using the latter, test the assumption of independence.

home	away							
	0	1	2	3	4	5	6	
0	32	33	9	14	3	0	1	
1	37	41	28	11	3	1	0	
2	27	25	29	7	2	1	0	
3	18	15	10	5	2	0	0	
4	9	6	3	0	0	1	0	
5	0	4	0	0	0	0	0	
6	0	1	2	0	0	0	0	

Table 1: Frequency of scores for EPL soccer matches in 2015

- (b) Fit the model characterized by Equation (E4.5.1) and answer the following questions:
- Using the fitted model, give the expected number of goals for a match between Manchester United (at home) against Liverpool.
 - Report and interpret the estimated home advantage $\hat{\delta}$.
 - Test whether the home advantage δ is significantly different from zero.
 - The asymptotic null distribution of the deviance statistic D is χ^2_{n-p-1} , but the latter is only valid when the number of observations in each group is large. In our analysis, there are only 38 matches in a given year at home/visiting for each team. We can instead approximate the null distribution of D using a simulation: specifically, we repeat the following steps $B = 10\,000$ times
 - generate new Poisson data from the fitted model
 - fit the Poisson regression specified by Equation (E4.5.1) on the simulated data
 - calculate the deviance statistic.
 Table 2 gives quantiles of the simulated null distribution of the deviance based on these $B = 10\,000$ simulations. Comment on the adequacy of the fit based on the deviance statistic and Table 2 and contrast with the conclusions obtained by using the asymptotic null distribution of the deviance.
- (c) Maher also suggested more complex models, including one in which the offensive and defensive strength of each team changed depending on whether they were at home or visiting another team, i.e.

$$Y_{ij1} \sim \text{Po}\{\exp(\alpha_i + \beta_j + \delta)\}, \quad Y_{ij2} \sim \text{Po}\{\exp(\gamma_j + \omega_i)\}, \quad i \neq j; i, j \in \{1, \dots, 20\} \quad (\text{E4.5.2})$$

Does Model E4.5.2 fit significantly better than Model E4.5.1?

- (d) Would a similar Poisson be adequate to model basketball scores? Justify your answer

4.6 **Bush vs Gore:** the 2000 US presidential election opposed Georges W. Bush (GOP) and Albert A. Gore (Democrat), as well as marginal third party candidates. The tipping state was Florida, worth 25 electors, which Bush won by a narrow margin of 537 votes. There have been many claims that the design of so-called butterfly ballots used in poor

level	2.5%	5%	10%	25%	50%	75%	90%	95%	97.5%
quantile	760.53	770.33	782.30	803.41	826.25	849.92	871.23	885.74	897.44

Table 2: Quantiles of the simulated deviance statistics based on the model of Maher (1982)

neighborhoods of Palm Beach county led to confusion among voters and that this deprived Gore of some thousands of votes that were instead assigned to a paleoconservative third-party candidate, Patrick Buchanan (Reform). Smith (2002) analysed the election results in Palm Beach country, in which a unusually high number of ballots (3407) were cast for Buchanan.

We are interested in building a model to predict the expected number of votes for Buchanan in Palm Beach county, based only on the information from other county votes. The buchanan data contains the following variables:

- `county`: name of county
- `popn`: population of the county in 1997.
- `white`: percentage of “white” (*sic*) in 1996 (per US Census definitions; people having origins in any of the original peoples of Europe, the Middle East, or North Africa).
- `black`: percentage of Black and African Americans in 1996 (origins in sub-saharian Africa).
- `hisp`: percentage of Hispanics in 1996.
- `geq65`: percentage of the population aged 65 and above based on 1996 and 1997 population estimates.
- `highsc`: percentage of the population with a high school degree (1990 Census data).
- `coll`: percentage of the population that are college graduates (1990 Census data).
- `income`: mean personal income in 1994.
- `buch`: total ballots cast for Pat Buchanan (Reform).
- `bush`: total ballots cast for Georges W. Bush (GOP).
- `gore`: total ballots cast for Al Gore (Democrat).
- `totmb`: total number of votes cast for the presidential election in each county, minus Buchanan votes.

(a) Calculate the total proportion of votes for Buchanan in Florida.

(b) Plot the percentage of votes obtained by Buchanan, $\text{buch}/(\text{buch}+\text{totmb})$, against $\ln(\text{popn})$ and comment.

Exclude the results of Palm Beach county for the rest of the question.

(c) We consider first a Poisson model for the percentage of votes for Buchanan, buch/totmb , as a function of `white`, $\ln(\text{hisp})$, `geq65`, `highsc`, $\ln(\text{coll})$, `income`.

- i. Explain why an offset is necessary in this case.
- ii. Why is `totmb` a better choice of denominator than `popn` for the rate? Explain.
- iii. Is the Poisson model appropriate? Justify your answer.
- iv. Explain why, if there is evidence of overdispersion, this means the binomial model is also inadequate.

Hint: what is the variance of the binomial distribution and how does it relate to the Poisson distribution?

(d) Use a negative binomial model with the same covariates to predict the expected number of Buchanan votes in Palm Beach county. Comment hence on the discrepancy between this forecast and the number of votes received in the election.